

# An Empirical Study of the Brute-Force Closest Pairs Algorithm

Jim Teresco  
jteresco@siena.edu

My Partner  
partner@fake.email

February 24, 2019

## Abstract

In this study, which is intended as an example of how to put together the writeup for an empirical study of algorithm performance, we present and analyze timings and operation counts for a straightforward brute-force closest pairs implementation. Points are generated randomly in a specified range, and the closest pair is computed and reported by a Java program.

## 1 Introduction

This empirical analysis study focuses on the brute-force closest pairs algorithm in two dimensions. We will consider a straightforward implementation of the algorithm in Java, instrumented to count basic operations and gather timings.

The remainder of this report is organized as follows. Section 2 describes the computing environment used. Section 3 describes the algorithm, our expectations based on the theoretical analysis, presents our results, and relates them back to the theory. We conclude with some additional discussion in Section 4.

## 2 Computing Environment

For this study, a Java program was developed that implements the algorithm of interest. Random inputs of varying sizes and other characteristics are generated, and the solution is computed.

All runs are using the following Java version:

```
java version "1.8.0_121"  
Java(TM) SE Runtime Environment (build 1.8.0_121-b13)  
Java HotSpot(TM) 64-Bit Server VM (build 25.121-b13, mixed mode)
```

The computing environment is a MacBook Pro (15-inch, 2016) with a 2.6 GHz Intel Core i7, 16 GB of 2133 MHz LPDDR3 memory, and running macOS Sierra Version 10.12.6. There are 4 cores, each with 256 KB L2 cache, and these cores share a 6 MB L3 cache.

### 3 Brute-Force Closest Pairs

The brute-force closest pairs algorithm as shown in Figure 1 was implemented in Java, and instrumented to count the number of times the distance between pairs of points is computed and to report the elapsed time for the main loops to execute.

**ALGORITHM** BRUTEFORCECLOSESTPOINTS( $P$ )

```
//Input: a set of points  $P[0..n - 1]$ 
 $d_{min} \leftarrow \infty$ 
for  $i \leftarrow 0..n - 2$  do
  for  $j \leftarrow i + 1..n - 1$  do
     $d \leftarrow \sqrt{(P[i].x - P[j].x)^2 + (P[i].y - P[j].y)^2}$ 
    if  $d < d_{min}$  then
       $d_{min} \leftarrow d$ 
       $index_1 \leftarrow i$ 
       $index_2 \leftarrow j$ 
return  $(index_1, index_2)$ 
```

Figure 1: The brute-force closest pairs algorithm, as implemented for this study.

The number of points for the runs is all powers of 2 from  $2^{10} = 1024$  to  $2^{19} = 524,288$ . Points for each run are generated randomly, within a range specified as a parameter to the program. Points are generated within the square with x- and y-coordinates no further than that range from the origin.

Each combination of the number of points and range of point positions is run a total of 5 times. The number of distance calculations is expected to be identical for each of the runs for a given number of points. Run times will vary, however, and the minimum time for a given combination of the number of points and range of point positions is chosen to minimize the effects from other computations that could be in process on the computer during the study.

#### 3.1 Expectations

The theoretical expectation is for both the number of distance calculations and elapsed time to be  $\Theta(n^2)$  in the best, average, and worst cases. There are no shortcuts out of the loops, which execute about  $\frac{n^2}{2}$  times.

## 3.2 Results

The raw results of the study are available in the file `timings.dat` in the GitHub repository. Each line, such as

```
BFCP 262144 64.0 92375.0 34359607296 (99936,203911) 4.069483000753879E-4
```

has 7 space-separated fields. The first field indicates the algorithm name, and is followed by the number of points, the range of coordinate values (each of  $x$  and  $y$  are within this distance of the origin), the elapsed time in milliseconds, the number of distance calculations, the indices of the two points that were found to be the closest pair, and the distance between those points. Of these, we are interested in the number of points, the elapsed time, and distance calculation count.

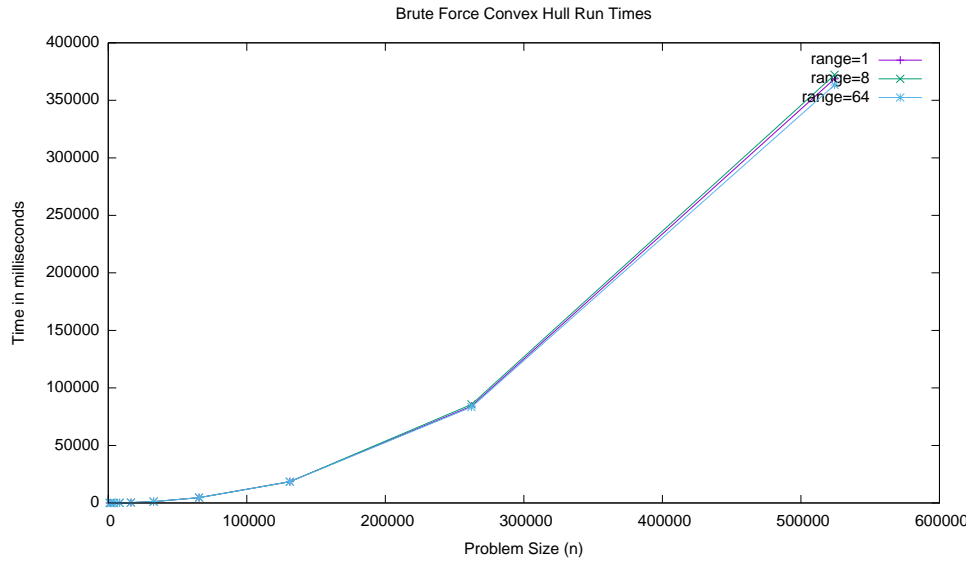


Figure 2: Times in milliseconds for the brute-force closest pairs to complete its computation for problem sizes from  $2^{10}$  to  $2^{19}$  points.

Figure 2 shows the times in milliseconds (minimum across 5 runs of each) and Figure 3 shows the number of computations of the distance between two points for this study. Unsurprisingly, the timings are very similar and the operation counts identical for the different ranges of point placements, so Table 1 shows the data as plotted in the figures but only for the range value of 1.0.

## 3.3 Discussion

Both the timings and distance calculation counts match well with the expected  $\Theta(n^2)$  behavior of this algorithm. The graphs in Figures 2 and 3 show the expected parabolic shape. The actual numbers in Table 1 align well. The number of distance calculations is exactly as predicted. Any variation there would have indicated an error in the algorithm or in the instrumentation that counted

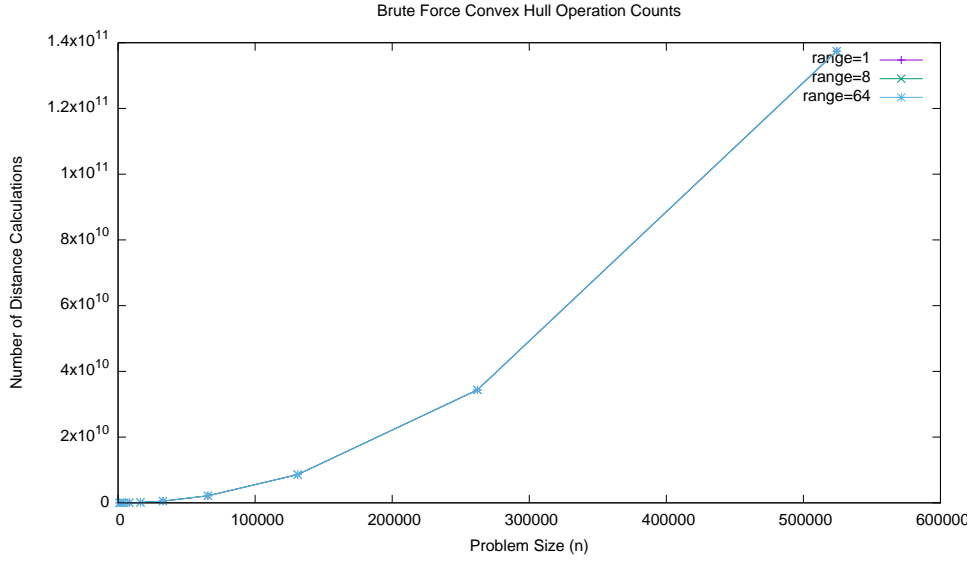


Figure 3: Number of distances between pairs of points computed the brute-force closest pairs to complete its computation for problem sizes from  $2^{10}$  to  $2^{19}$  points.

the operations. For the timings, small problem sizes are subject to errors due to the accuracy of the millisecond timer over short timespans. However, when we look at the larger problem sizes, we see exactly what we would expect: when we double the problem size, the time taken to compute the solution increased by a factor of four.

## 4 Conclusions

In this simple study, we found that all of our results matched the expected behavior from the theory. In this particular algorithm, since there are no asymptotic differences among the best-, average-,

Problem Size	Time (ms)	Distance Calculations
1024	11.0	523776
2048	15.0	2096128
4096	35.0	8386560
8192	81.0	33550336
16384	298.0	134209536
32768	1171.0	536854528
65536	4607.0	2147450880
131072	18496.0	8589869056
262144	84123.0	34359607296
524288	368682.0	137438691328

Table 1: Actual times and operation counts used to create the plots in Figure 2 and Figure 3.

and worst-case behavior, this should be unsurprising. The small variations in run time can perhaps be explained by how soon the algorithm finds the closest pair (or at least a very close pair). The more times that a new closest pair is found, the more often the body of the `if` statement following the distance calculation will need to execute. There is also the issue that the computer used for the study is a laptop system that had many processes such as web browsers with many tabs and terminal windows in regular interactive use during the timing studies. This is partially accounted for by taking the minimum time across 5 runs, but as each run has its own set of random points generated, the previously-mentioned effect from the different numbers of times the body of the `if` statement is executed will potentially skew the results more toward a best case behavior. Any such variations are likely minor, and do not decrease the confidence in our conclusions that the timing results match closely with theoretical expectations.